

Practical 3 Solutions

Jumping Rivers

Predictive Analytics: practical 3

The OJ data set

The OJ data set from the ISLR package contains information on which of two brands of orange juice customers purchased¹ and can be loaded using

```
data(OJ, package = "ISLR")
```

After loading the caret and jrPred package

```
library("caret")
library("jrPred")
```

make an initial examination of the relationships between each of the predictors and the response²

```
par(mfrow = c(4, 5), mar = c(4, 0.5, 0.5, 0.5))
plot(Purchase ~ ., data = OJ)
```

Initial model building using logistic regression

- To begin, create a logistic regression model that takes into consideration the prices of the two brands of orange juice, PriceCH and PriceMM. Hint: Use the `train` function, with `method = 'glm'`. Look at the help page for the data set to understand what these variables represent.

```
m1 = train(Purchase ~ PriceCH + PriceMM,
           data = OJ, method = "glm")
```

- What proportion of purchases does this model get right?

```
getTrainPerf(m1)
```

```
##      TrainAccuracy TrainKappa method
## 1      0.6181051 0.07974531      glm
```

- How does this compare to if we used no model?

```
# with no model we essentially predict according to
# proportion of observations in data
```

```
# work out proportions
probs = table(OJ$Purchase)/nrow(OJ)
# sample using proportions
preds = sample(levels(OJ$Purchase), prob = probs)
# work out correct proportion
mean(preds != OJ$Purchase)
```

```
## [1] 0.5009346
```

¹The response variable is `Purchase`.

²Use the `plot` function with a model formula or the `pairs` function.

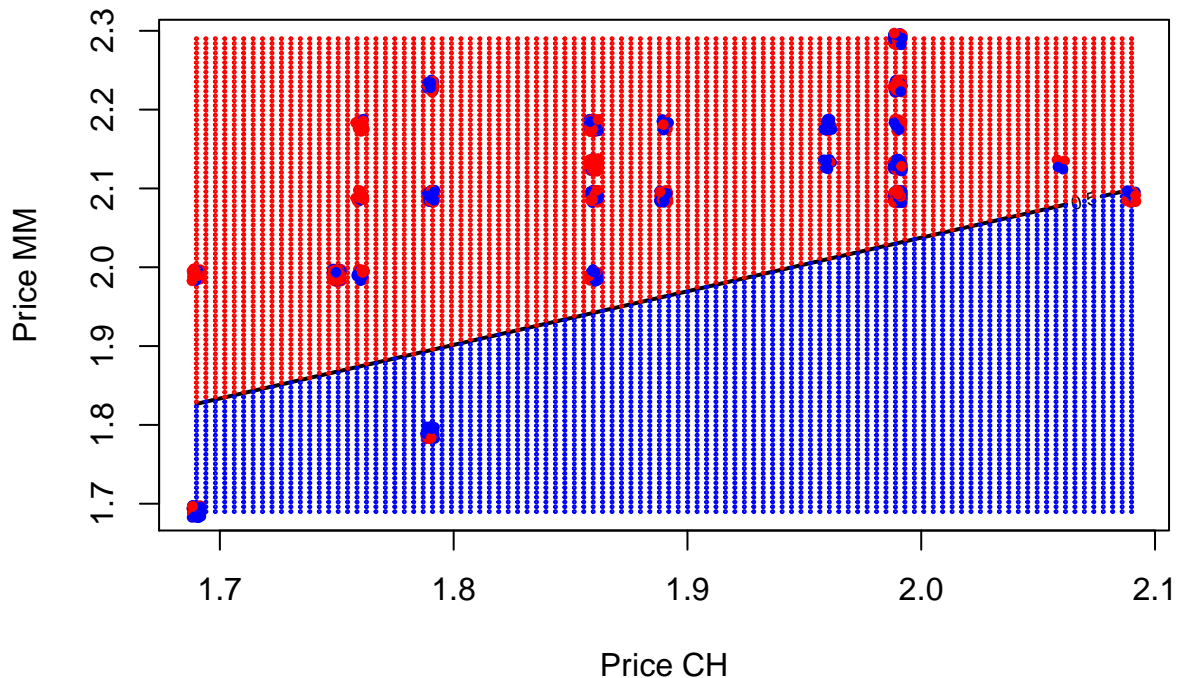


Figure 1: Examining the decision boundary for orange juice brand purchases by price.

- Use your model to predict if a customer will buy CH or MM if the price of CH and MM is 2.3 and 2.4 respectively

```
predict(m1, newdata = data.frame(PriceCH = 2.3, PriceMM = 2.4))
```

```
## [1] CH
## Levels: CH MM
```

Visualising the boundary

The `jrPred` package contains following code produces a plot of the decision boundary as seen in figure 1.

```
boundary_plot(m1, OJ$PriceCH, OJ$PriceMM, OJ$Purchase,
              xlab="Price CH", ylab="Price MM")
```

Run the boundary code above, and make sure you get a similar plot.

- What happens if we add an interaction term? How does the boundary change?

```
# We now have a curved decision boundary.
# There are two regions of where we would predict MM, bottom left, and a tiny one up in the top right.
```

- Try adding polynomial terms.

Using all of the predictors

- Instead of just using 2 predictors we want to use all of them. However, we have a few problems to tackle first. A few of our predictors are linear combinations of the others. This leads to what is called rank-deficiency problems. For instance, if you run the following model you'll realise there are a few NAs.

```
mLM = train(Purchase ~ ., data = OJ, method = "glm")
```

Take the predictor PriceDiff. It is impossible to estimate it's coefficient as it is a linear combination of PriceCH and PriceMM i.e. $\text{PriceDiff} = \text{PriceCH} - \text{PriceMM}$. In this particularly data set, there are quite a few linear combinations. We can find them using the `findLinearCombos()` and `model.matrix()` functions

```
remove = findLinearCombos(model.matrix(Purchase ~ ., data = OJ))
```

The output list has a component called `remove` suggesting which variables should be removed to get rid of linear combinations

```
(badvar = colnames(OJ)[remove$remove])
```

```
## [1] "SalePriceMM" "SalePriceCH" "PriceDiff" "ListPriceDiff"
## [5] "STORE"
```

We can then remove these variable from the data

```
OJsub = OJ[, -remove$remove]
```

- Use the new OJsub data set to model Purchase using all of the predictors. How accurate is the model?

```
mLM = train(Purchase ~ ., data = OJsub, method = "glm")
getTrainPerf(mLM)
```

```
## TrainAccuracy TrainKappa method
## 1 0.8255619 0.6309913 glm
```

- What are the values of sensitivity and specificity?

```
## could use confusionMatrix
(cmLM = confusionMatrix(predict(mLM,OJsub),OJsub$Purchase))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  CH  MM
##           CH 577 100
##           MM  76 317
##
##           Accuracy : 0.8355
##           95% CI : (0.8119, 0.8572)
##           No Information Rate : 0.6103
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.6506
##           Mcnemar's Test P-Value : 0.08297
##
##           Sensitivity : 0.8836
##           Specificity : 0.7602
##           Pos Pred Value : 0.8523
##           Neg Pred Value : 0.8066
##           Prevalence : 0.6103
##           Detection Rate : 0.5393
##           Detection Prevalence : 0.6327
##           Balanced Accuracy : 0.8219
##
##           'Positive' Class : CH
##
```

```
# or
sensitivity(predict(mLM,OJsub),OJsub$Purchase)
```

```
## [1] 0.8836141
```

```
specificity(predict(mLM,OJsub),OJsub$Purchase)
```

```
## [1] 0.7601918
```

- What does this mean?

#The model is fairly good at picking up both positive events, person buys CH, and negative events, MM.

K nearest neighbours

- Try fitting models using the K nearest neighbours algorithm. To begin with, just have two covariates and use the `boundary_plot` function to visualise the results.

```
mKNN = train(Purchase~., data = OJsub, method = "knn")
```

- How do they compare in accuracy, sensitivity and specificity?

```
cmKNN = confusionMatrix(predict(mKNN,OJsub),OJsub$Purchase)
(info = data.frame(Model = c("logistic","knn"),
  Accuracy = c(cmLM$overall["Accuracy"],
    cmKNN$overall["Accuracy"]),
  Sensitivity = c(cmLM$byClass["Sensitivity"],
    cmKNN$byClass["Sensitivity"]),
  Specificity = c(cmLM$byClass["Specificity"],
    cmKNN$byClass["Specificity"])))
```

```
##      Model  Accuracy Sensitivity Specificity
## 1 logistic 0.8355140   0.8836141   0.7601918
## 2      knn 0.8065421   0.8928025   0.6714628
```

- How does varying the number of nearest neighbours in a KNN affect the model fit?

Accuracy increases at first with knn before then getting worse after a peak value of 9.

```
(mKNN2 = train(Purchase~., data = OJsub, method = "knn",
  tuneGrid = data.frame(k = 1:30)))
```

```
## k-Nearest Neighbors
##
## 1070 samples
## 12 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1070, 1070, 1070, 1070, 1070, ...
## Resampling results across tuning parameters:
##
##  k    Accuracy    Kappa
##  1  0.6981705  0.3643238
##  2  0.6873374  0.3411474
##  3  0.6911386  0.3492862
##  4  0.6944823  0.3546660
```

```
##    5  0.7028483  0.3710620
##    6  0.7098169  0.3835293
##    7  0.7130895  0.3879166
##    8  0.7076557  0.3748091
##    9  0.7080973  0.3744772
##   10  0.7074590  0.3709005
##   11  0.7056465  0.3671281
##   12  0.7012401  0.3563953
##   13  0.6970988  0.3473017
##   14  0.6959412  0.3433170
##   15  0.6916124  0.3327993
##   16  0.6886968  0.3260372
##   17  0.6896934  0.3270047
##   18  0.6835671  0.3147028
##   19  0.6843347  0.3153751
##   20  0.6802114  0.3057056
##   21  0.6796077  0.3040183
##   22  0.6757461  0.2958511
##   23  0.6741128  0.2922900
##   24  0.6722414  0.2884721
##   25  0.6721267  0.2873917
##   26  0.6731688  0.2889079
##   27  0.6729676  0.2882316
##   28  0.6723921  0.2861108
##   29  0.6739748  0.2892669
##   30  0.6763646  0.2950767
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 7.
```

The KNN algorithm described in the notes can also be used for regression problems. In this case the predicted response is the mean of the k nearest neighbours.

- Try fitting the KNN model for the regression problem in practical 1.

```
library("jrPred")
data(FuelEconomy, package = "AppliedPredictiveModeling")
regKNN = train(FE~., data = cars2010, method = "knn")
regLM = train(FE~., data = cars2010, method = "lm")
getTrainPerf(regKNN)
```

```
##    TrainRMSE TrainRsquared TrainMAE method
## 1  3.526938      0.7872847 2.422289      knn
```

```
getTrainPerf(regLM)
```

```
##    TrainRMSE TrainRsquared TrainMAE method
## 1  3.787227      0.75811 2.546495      lm
```

- How does this compare to the linear regression models?

```
# The KNN regression model is not as good as the linear model, only just
```

Resampling methods

- Fit a KNN regression model to the `cars2010` data set with `FE` as the response.

```
data(FuelEconomy, package = "AppliedPredictiveModeling")
```

```
mKNN = train(FE ~ ., method = "knn", data = cars2010)
```

- Estimate test error using 10-fold cross validation

```
# set the train control object
tc10fold = trainControl(method = "cv", number = 10)
# fit the model using this train control object
mKNN10 = train(FE ~ ., method = "knn", data = cars2010,
               trControl = tc10fold)
getTrainPerf(mKNN10)
```

```
##   TrainRMSE TrainRsquared TrainMAE method
## 1   3.337385      0.8053001 2.300853   knn
```

- Again using 10 fold CV, estimate the performance of the k nearest neighbours algorithm for different values of k .

```
mKNNcv10 = train(FE ~ ., method = "knn", data = cars2010,
                 trControl = tc10fold, tuneGrid = data.frame(k = 2:20))
```

- Which model is chosen as the best?

```
mKNNcv10$bestTune
```

```
##    k
## 1  2
```

- Create new `trainControl` objects to specify the use of 5 fold and 15 fold cross validation to estimate test RMSE.

```
tc5fold = trainControl(method = "cv", number = 5)
tc15fold = trainControl(method = "cv", number = 15)
```

- Go through the same training procedure attempting to find the best KNN model.

```
mKNNcv5 = train(FE ~ ., data = cars2010, method = "knn",
                trControl = tc5fold, tuneGrid = data.frame(k = 2:20))

mKNNcv15 = train(FE ~ ., data = cars2010, method = "knn",
                 trControl = tc15fold, tuneGrid = data.frame(k = 2:20))
mKNNcv5$bestTune
```

```
##    k
## 3  4
```

```
mKNNcv15$bestTune
```

```
##    k
## 1  2
```

An example with more than two classes

The `Glass` data set in the `mlbench` package is a data frame containing examples of the chemical analysis of 7 different types of glass. The goal is to be able to predict which category glass falls into based on the values of the 9 predictors.

```
data(Glass, package = "mlbench")
```

A logistic regression model is typically not suitable for more than 2 classes, so try fitting a k nearest neighbour model. Use k-fold cross validation if you want to. What proportion of predictions does your model get correct?

```
tc = trainControl(method = "cv", number = 10)
model = train(Type ~ ., data = Glass, trControl = tc, method = "knn")
getTrainPerf(model)
```

```
##      TrainAccuracy TrainKappa method
## 1          0.686118  0.5629221     knn
```